

How CostGuard Actually Works

A plain-language walkthrough — no prior knowledge of "tokens" required.

1 What's a token?

AI reads and writes in small chunks called tokens (~4 letters each). Every word in and out gets counted.

2 Why does it cost money?

AI companies charge per token — like a phone bill charges by the minute. Tiny amounts, but they add up fast.

3 6 stages per task

Spec → Plan → Implement → Test → Review → Debug. Each stage is its own AI conversation, its own cost.

ONE SENTENCE It reads a bill and a budget, and tells you where you overspent. It doesn't build or stop anything — it just flags problems, like a banking app flagging overspending.

EXAMPLE One Real Task, All 6 Stages — Task T-002 (an easy task, difficulty 1 of 5)

STAGE	AI MODEL USED	COST
Spec	Cheap model	< 1¢
Plan	Expensive model ⚠	< 1¢ (0.75¢)
Implement	Expensive model	1.4¢
Test	Expensive model	0.9¢
Review	Expensive model	3.2¢
Debug	Expensive model	1.7¢
Total for this one task		≈ 8¢

THE 3 QUESTIONS What CostGuard Checks, Every Time It Runs

QUESTION 1 — GOVERN

Did we go over budget anywhere?

YES — flagged

All "Review" stages, added up across every task this month, hit \$48.63 against a \$35 budget.

QUESTION 2 — MEASURE

Did any ONE task cost way more than it should have?

NO — all clear

T-002's 8¢ is normal for an easy task. Nothing in this batch looks abnormal for its difficulty.

QUESTION 3 — MANAGE

Did we use the expensive model when the cheap one would've worked?

YES — 75¢ wasted

T-002's Plan stage is 1 of 24 cases where easy work used the expensive model by accident.

BOTTOM LINE, THIS RUN

\$138.29 total spend across 50 tasks · **1 phase over budget** (Review, by \$13.63) · **0 tasks abnormally expensive** · **\$0.75 recoverable** (92% waste) from model-routing mistakes — small in dollars on 50 tasks, not small at a real company's monthly volume.